

Shubham Kumar Gupta

+91-9376091240 | kshubham04907@gmail.com | shubham.cv | linkedin.com/in/shubhamgupta04907 | github.com/kshubham090

FOCUS

AI engineer who builds and ships agentic infrastructure: eval harnesses, memory-backed gateways, and replay tooling for debugging agent behavior. I calibrate before I trust a system, which is why the eval work exists in the first place. No hands-on model training yet. What I have is real evaluation depth, a completed voice pipeline that replaced a vendor API at a fraction of the cost, and a research habit that carries into how I build.

EXPERIENCE

WINNIIO / LifeAtlas

Sweden (Remote)

AI Engineering Intern

May 2026 – Jul 2026

- Led a six-person team rebuilding a production voice-interview pipeline in-house after Retell AI's per-call cost became the real constraint on scaling, cutting cost roughly 20x with zero downtime during the swap.
- Original prototype ran voice in, Claude-based profile extraction, and ranked matching via pgvector and Voyage rerank on Retell AI and Twilio.
- Also touched the ZeroClaw auth proxy and edge-native AI work on the LifeAtlas platform.

Stakrid Logistics Private Limited

Noida, India

Founder, Lead Engineer

Jan 2025 – Jan 2026

- Built a last-mile logistics platform solo: 40+ REST endpoints (Java, Spring Boot) on GCP with CI/CD, load balancing, and monitoring.
- Query optimization and connection pooling cut manual processing 80% and API latency from 800ms to under 200ms.
- Integrated Supabase, payment gateways, and SMS APIs with retry logic. Wound down deliberately once it had run its course.

PROJECTS

Agent Reliability Stack (Python, Claude API, FastAPI, Neo4j, pgvector, Voyage AI, Redis, GitHub Actions) Jun 2026 – Present

- Eval harness with pluggable scorers (exact, regex, embedding, LLM-as-judge) and LCS-based trajectory scoring on tool-call sequences, thread-pooled with failure isolation.
- Validated the LLM-judge scorer against 60 human-labeled cases before trusting it: exact-match 0.867, embedding similarity 0.984, calibrated judge 1.000. Wired into GitHub Actions as a baseline-diffed regression gate.
- LLM gateway backed by a Neo4j memory graph. One Voyage-3.5 embedding call drives semantic cache lookup, graph seeding, and write-back together; cache misses trigger a 1–2 hop graph walk that injects ranked context.
- Auto-fallback to a secondary model tier on 429/timeout/5xx, with per-call cost attribution by user and feature.

Relogged: Sandboxed Tool-Call Replay Debugger (Python, LangGraph, Postgres)

Jul 2026

- Zero-code-change recorder that monkeypatches `CompiledStateGraph` and `ToolNode` via a scoped `@record()` contextvar. Pure replay reconstructs a run from the database alone with zero live model, tool, or graph invocation.

RESEARCH

Auditing Evaluation Leakage in F1 Race-Strategy ML (doi.org/10.5281/zenodo.21862868)

Jul 2026

- Found a prior lap-time pipeline's $R^2=0.994$ was leakage from row-wise splits plus one-hot track features; a corrected track-aware protocol restored an honest $R^2=0.991$.

Vyaskosh / VAYAS: Age-Stratified ASR Fairness Audit (*protocol design phase*)

2026 – Present

- Portable audit protocol for elderly-speaker ASR evaluation, following the Gender Shades build-benchmark-audit methodology across Indo-Aryan and Dravidian languages.

Chakra47: Autonomous OS for Physical AI (chakra47.com)

Dec 2025 – May 2026

- Four-layer LangGraph multi-agent swarm with neuro-symbolic governance (YOLOv8 + Open Policy Agent), presented at India AI Impact Summit 2026 and AICraft 3.1.

EDUCATION

Amity University Noida

Noida, India

B.Tech, Computer Science, AI & ML Specialization

Aug 2023 – Jul 2027

- Self-study: Stanford CME295 (Transformers & LLMs); Andrej Karpathy's *Neural Networks: Zero to Hero*.
- Certifications: Machine Learning Specialization (Andrew Ng, Stanford/Coursera) · NCA-AIIO: AI Infrastructure & Operations Associate (NVIDIA, in progress) · Java Spring Framework 6 with Spring Boot 3 (Udemy) · Build and Secure Networks in Google Cloud (Google Cloud Skills Boost).

SKILLS

Languages: Python, Java, SQL, Bash

DSA: Arrays, Linked Lists, Trees, Graphs, Dynamic Programming, Recursion, Sorting, Sliding Window, Two Pointers

Agentic & AI Infra: LangGraph, LangChain, Claude API, RAG, pgvector, Voyage rerank, Open Policy Agent, sentence-transformers

ML & CV: PyTorch, TensorFlow, CNNs, RNNs, LSTMs, Transformers, YOLOv8, OpenCV, Hugging Face

Backend: Spring Boot, REST API Design, FastAPI, JDBC, Spring Data JPA, PostgreSQL, Neo4j, SQLAlchemy

Cloud & DevOps: Google Cloud Platform, Docker, CI/CD, GitHub Actions, Supabase, Firebase